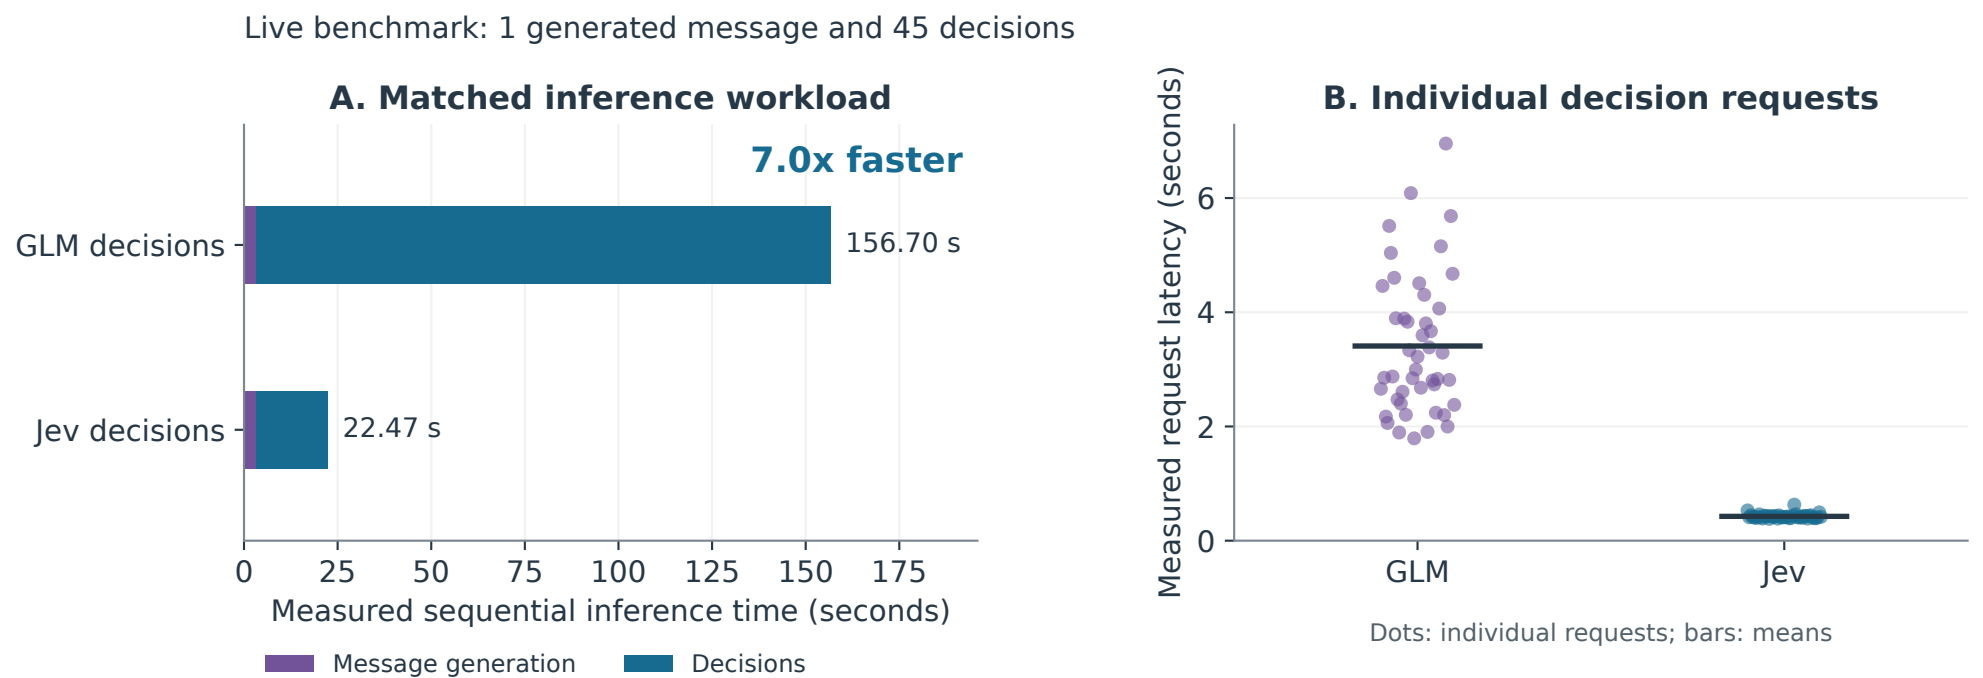
Live benchmark: 1 generated message and 45 decisions
A. Matched inference workload
7.0x faster
GLM decisions
Jev decisions
156.70 s
22.47 s
0
25
50
75
100
125
150
175
Measured sequential inference time (seconds)
Message generation
Decisions
B. Individual decision requests
Measured request latency (seconds)
0
2
4
6
GLM
Jev
Dots: individual requests; bars: means